

# Fahim Tahmid Chowdhury

Mobile: +1(786)-406/2617  
Skype: fahim.tahmid.chowdhury  
E-mail: fchowdhu@cs.fsu.edu  
Website: <http://ww2.cs.fsu.edu/~fchowdhu>  
LinkedIn: <http://www.linkedin.com/in/fahimtahmidchowdhury>

## EDUCATION

**Florida State University**, Tallahassee, Florida

Ph.D. candidate, Computer Science, currently enrolled, *CGPA 4.00/4.00* (Intended grad: **January 2022**)

**Bangladesh University of Engineering and Technology**, Dhaka, Bangladesh

BSc., Computer Science and Engineering, February 2013, *CGPA 3.54/4.00*

## RESEARCH INTERESTS

- **High Performance Computing (HPC) Systems:** Domain-specific Systems Design, HPC I/O Optimization, Heterogeneous Storage Stack, HPC File Systems, HPC Workflow, Performance Analysis
- **Artificial Intelligence:** Deep Learning(DL) at Scale, Reinforcement Learning

## RESEARCH EXPERIENCE

**Department of Computer Science, Florida State University**

**Graduate Research Assistant**

**August 2017 - Present**

- Ph.D. student researcher at *Computer Architecture and SysTems Research Lab (CASTL)* supervised by *Professor Dr. Weikuan Yu*, specializing in domain-specific distributed systems design

**Center for Applied Scientific Computing (CASC)**, Lawrence Livermore National Laboratory (LLNL)  
**Student Intern**

**May 2019 - August 2019**

- Worked on optimizing I/O strategies in HPC application workflows like Cancer Moonshot Pilot 2 in the *Data Analysis Group* at CASC. Achieved 84.7% latency improvement by using burst buffers on Lassen.

**National Energy Research Scientific Computing Center (NERSC)**, Lawrence Berkeley National Laboratory (LBNL), Berkeley, California

**Student Assistant**(Summer intern)

**May 2018 - August 2018**

- Worked in the *Data Analytics and Services* group at NERSC. Analyzed scalable data pipeline for distributed DL atop TensorFlow and Horovod. Determined I/O bottleneck of upto 11.04% in DL training time.

**NERSC, LBNL**

**LBNL Affiliate**

**August 2018 - August 2019**

- Enhanced the summer internship project on determining I/O bottlenecks in distributed DL applications.

## RESEARCH PROJECTS

- **HPC Workflow I/O Optimization:** Built an Emulator during internship at CASC to analyze different HPC I/O patterns, e.g., Deep Learning Training I/O, Checkpoint/Restart, Producer-Consumer, etc. Developing a middleware to implement optimization strategies for application workflows.
- **BeeGFS Performance Evaluation:** Published a research paper on the performance evaluation of *BeeGFS* parallel cluster file system using IOR and MDTest, and deep learning applications atop TensorFlow, Horovod and LBANN. Currently, working on analyzing the fitness of *BeeGFS On Demand (BeeOND)* as an ephemeral burst buffer file system.
- **Scalable Data Pipeline for Distributed Deep Learning:** Analyzed and profiled I/O behavior posed by cutting-edge deep learning applications at scale by using a *logging framework* developed during internship at NERSC. Pinpointed I/O bottlenecks caused by metadata overhead in deep learning training.

## PUBLICATIONS

- **F. Chowdhury**, F. Di Natale, A. Moody, E. Gonsiorowski, K. Mohror, and W. Yu. "Understanding I/O Behavior in Scientific Workflows on High Performance Computing Systems," in Proceedings of the *International Conference on High Performance Computing, Networking, Storage and Analysis 2019 (SC19)*, *Regular Poster*, Nov. 2019.
- **F. Chowdhury**, Y. Zhu, T. Heer, S. Paredes, A. Moody, R. Goldstone, K. Mohror, and W. Yu, "I/O Characterization and Performance Evaluation of BeeGFS for Deep Learning," in Proceedings of the *48th International Conference on Parallel Processing (ICPP 2019)*, *Research Paper*, Aug. 2019.
- **F. Chowdhury**, J. Liu, Q. Koziol, T. Kurth, S. Farrell, S. Byna, Prabhat, and W. Yu, "Initial Characterization of I/O in Large-Scale Deep Learning Applications," in *SC18, 3RD Joint International Workshop on Parallel Data Storage & Data Intensive Scalable Computing Systems (PDSW-DISCS 2018)*, *Work-in-progress (WIP) Abstract*, Nov. 2018.
- Y. Zhu, **F. Chowdhury**, H. Fu, A. Moody, K. Mohror, K. Sato, and W. Yu, "Entropy-Aware I/O Pipelining for Large-Scale Deep Learning on HPC Systems," in *IEEE International Symposium on the Modeling, Analysis, and Simulation of Computer and Telecommunication Systems (MASCOTS 2018)*, *Research Paper*, Sep. 2018.

## TECHNICAL SKILLS

- Programming Languages: **C/C++**, **Python**, C#, MATLAB, Java, Javascript
- Libraries: **MPI**, **HDF5**, BSD sockets, WinSock, OpenGL, Boost, Windows API, Google Test
- Frameworks: **TensorFlow**, **Horovod**, **LBANN**, Qt Framework, MFC, .NET Framework
- Distributed File Systems: **BeeGFS**, Lustre, BurstFS, UnifyCR